

# Vaibhav Gogte

2260 Hayward St., Ann Arbor, MI 48109 | [vgogte@umich.edu](mailto:vgogte@umich.edu) | +1 (352) 870-7970 | <http://web.eecs.umich.edu/~vgogte>

---

## Research Interests

My current research focuses on designing runtime systems and hardware architecture for emerging byte-addressable persistent memories. Specifically, I have designed memory models in high-level languages and processor architectures that allow easier persistent memory programming, and built software systems that use persistent memories as faster storage.

## Education

- **Ph.D. Candidate, Computer Science and Engineering** Sept 2014 – Dec 2019 (expected)  
**M.S. Computer Science and Engineering** Sept 2014 – Apr 2016  
*University of Michigan - Ann Arbor* (GPA: 4.0/4.0)  
Thesis: Runtime Systems for Persistent Memories  
Advisor: Prof. Thomas F. Wenisch
- **B.E. (Hons.), Electrical & Electronics Engineering** Aug 2007 – May 2011  
*Birla Institute of Technology & Science - Pilani, India* (GPA: 8.92/10.0)

## Experience

**Graduate Student Research Assistant, University of Michigan, Ann Arbor** Sept 2014 – Present

- **Strand Persistency:** In this work, I proposed hardware primitives to minimally constrain ordering on persistent memory operations. I designed compiler mechanisms that map persistency semantics in high-level languages to the hardware primitives. This work achieves up to  $1.97\times$  speedup over existing hardware ordering mechanisms.
- **Persistency for Synchronization-Free Regions:** In this work, I proposed persistency semantics that guarantee failure atomicity of synchronization-free regions, program regions delimited by synchronization operations. Our approach provides clear semantics for post-failure state by extending sequentially consistent guarantees to the recovery code.
- **Language-Level Persistency:** This project defines an acquire-release persistency model as an extension to the C++11 memory model. The persistency model relaxes the ordering constraints while accessing recoverable data structures in persistent memory-enabled systems.
- **OS-based Wear Management for Persistent Memories:** Emerging persistent memories suffer from a limited write endurance. I designed an application-transparent wear-management technique in Linux kernel that detects the disparity in the write accesses and performs targeted page migrations to manage wear of the persistent memory.
- **Hardware Accelerator for Regular Expressions:** In this work, I designed a stall-free hardware accelerator for matching regular expressions at a scan rate offered by modern DDRs. I built a compiler toolchain in C++ that transforms the regular expressions to the binary used by our hardware design.
- **Programmable and Energy-Efficient 3D Convolution Engine:** I developed a hardware accelerator that improves data locality when fetching 3D images from the main memory. The convolution engine overlaps the processing of locally cached image pixels with the fetching of a new set of pixel data from the main memory.

**Performance and Capacity Research Intern, Facebook, Menlo Park** May 2019 – Aug 2019

- **Persistent Memory Usecases for Facebook Applications:** In this work, I designed persistent memory-aware block cache for RocksDB, a key-value store application. I evaluated several RocksDB usecases in Facebook applications to show throughput and latency improvement due to the persistent memory-aware cache.

**Research Intern, Microsoft Research, Redmond** May 2018 – Aug 2018, May 2017 – Aug 2017

- **Low-cost and Predictable Storage Management for Persistent Memories:** The access latency of persistent memories, storage features and modern networks are a few microseconds. In this work, I proposed low-cost and predictable storage by offloading their management to the ARM-based commodity SoC device attached to the host over PCIe.
- **Load-balancing in Catapult architecture:** I extended the routing and transport layer in the Catapult architecture to dynamically balance query load in the FPGA-enabled servers. The proposed mechanism monitors the response latency and manages the load offered by the FPGAs.

**Senior Design Engineer, Texas Instruments, Bangalore** Jul 2011 – Jul 2014

- **Design of microprocessor subsystems based on ARM processors:** I designed power management, protocol converters, and interrupt controllers for dual-core microprocessor subsystems based on ARM Cortex A9, A7, and M4 processors.

## Teaching Experience

**Graduate Student Instructor**, *University of Michigan, Ann Arbor*

- Parallel Computer Architecture (EECS 570) Jan 2017 – May 2017
- Introduction to Logic Design (EECS 270) Jan 2015 – May 2015

## Peer-Reviewed Publications

- **StrandWeaver: Relaxed Persist Ordering Using Strand Persistency** (Under review)  
V. Gogte, W. Wang, S. Diestelhorst, P. M. Chen, S. Narayanasamy, and T. F. Wenisch.
- **Language Support for Memory Persistency**  
A. Kolli, V. Gogte, A. Saidi, S. Diestelhorst, W. Wang, P. M. Chen, S. Narayanasamy, and T. F. Wenisch.  
Top Picks of the 2019 Computer Architecture Conferences (**Top Picks**). Jun 2019.
- **Software Wear Management for Persistent Memories**  
V. Gogte, W. Wang, S. Diestelhorst, A. Kolli, P. M. Chen, S. Narayanasamy, and T. F. Wenisch.  
17th USENIX Conference on File and Storage Technologies (**FAST**). Feb 2019.
- **Persistency for Synchronization-Free Regions**  
V. Gogte, W. Wang, S. Diestelhorst, S. Narayanasamy, P. M. Chen, and T. F. Wenisch.  
39th ACM SIGPLAN Conference on Programming Language Design and Implementation (**PLDI**). Apr 2018.
- **Language-Level Persistency**  
A. Kolli, V. Gogte, A. Saidi, S. Diestelhorst, P. M. Chen, S. Narayanasamy, and T. F. Wenisch.  
44th Annual International Symposium on Computer Architecture (**ISCA**). Jun 2017.
- **HARE: Hardware Accelerator for Regular Expressions**  
V. Gogte, A. Kolli, M. J. Cafarella, L. D'Antoni, T. F. Wenisch.  
49th International Symposium on Microarchitecture (**MICRO**). Oct 2016.

## Workshops

- **Strand Persistency**  
V. Gogte, W. Wang, S. Diestelhorst, P. M. Chen, S. Narayanasamy, and T. F. Wenisch.  
10th Annual Non-Volatile Memories Workshop (**NVMW**). Mar 2019.
- **Failure-Atomic Synchronization-Free Regions**  
V. Gogte, W. Wang, S. Diestelhorst, S. Narayanasamy, P. M. Chen, and T. F. Wenisch.  
9th Annual Non-Volatile Memories Workshop (**NVMW**). Mar 2018.
- **TARP: Translating Acquire-Release Persistency**  
A. Kolli, V. Gogte, A. Saidi, S. Diestelhorst, P. M. Chen, S. Narayanasamy, and T. F. Wenisch.  
8th Annual Non-Volatile Memories Workshop (**NVMW**). Mar 2017.

## Patents

- **Instruction ordering**  
P. M. Chen, S. Diestelhorst, V. Gogte, S. Narayanasamy, W. Wang, T. F. Wenisch. (*Patent Pending*)
- **Detecting at least one predetermined pattern in stream of symbols**  
M. J. Cafarella, V. Gogte, T. F. Wenisch. (*Patent Pending*)

## Press

- Persistency for Synchronization-Free Regions, *ARM Research*, 2018. [Link](#).
- Honorable mentions for service as EECS 570 graduate student instructor, 2017. [Link](#).
- Baking Specialization into Hardware Cools CPU Concerns, *Next Platform*, 2016. [Link](#).

## Service

- **Shadow Program Committee Member** for EuroSys'19
- **Reviewer** for IEEE CAL'19, TOMPECS'19, TACO'18 and TODAES'16 journals

## Technical Skills

- **Languages:** C, C++, Verilog HDL
- **Software and Design Tools:** LLVM, Murphi, DynamoRio, Git, Latex
- **Architectural simulators:** Gem5, Booksim